

Drive_For_Show_Putt_For_Dough

May 3, 2017

0.0.1 Steve DeLano, CSCI E-7 Final Project, May 3, 2017

0.0.2 Drive for Show, Putt for Dough?

0.0.3 This common expression in golf refers to the idea that the first shot on each hole (the drive) goes the longest distance (290 yards for the average professional golfer), but the shortest shot on each hole (the putt -- which is often 20 feet or less) is what separates the best golfers from the rest of the pack.

0.0.4 The saying is so common that the concept is almost universally accepted. But, is it true?

0.0.5 Well, let's put Python to work for the CSCI E-7 final project and find out.

0.0.6 Here is an outline of the key steps in the process:

1. Locate dataset of golf statistics
2. Identify key libraries and import them
3. Read and explore data; clean if necessary
4. Analyze the data
5. Create key data visualizations to tell a story

0.1 Spoiler Alert!!!

0.1.1 On the PGA Tour, driving effectiveness has a higher correlation with winning money than the correlation between putting and winning money.

0.1.2 BUT ... neither is as important as being effective from "tee to green". That is, the players who make it from the tee to the green in the fewest shots are the ones who make the most money.

0.1.3 Let's go step by step through the process and see the data that brings this conclusion to life.

0.1.4 1. Locate dataset of golf statistics.

The source of the data is pgatour.com/stats. My son, Jimmy, put together a program to scrape key statistics from the website. He used Java and downloaded data that he converted into a file

called 'golf.csv'. Jimmy is a 5 handicap golfer (that means he's really good). I'm a 25 handicap golfer (that means I'm really bad). Fortunately, Jimmy is also good at writing webscraping code, and since Lena has not given her presentation on webscraping yet, I decided it was ok to enlist Jimmy's skills in putting together the dataset. Thanks, Jimmy!

0.1.5 2. Identify key libraries and import them

```
In [1]: import solution
```

```
-----  
  
ImportError                                Traceback (most recent call last)  
  
  <ipython-input-1-5fc6bd43c1ae> in <module>()  
----> 1 import solution  
  
ImportError: No module named 'solution'
```

0.1.6 Nenad would say: "Come on guys. You have to try harder than that!"

0.1.7 Ok, back to the drawing board.

```
In [2]: import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt  
import scipy.stats as stats  
import warnings  
from bokeh.plotting import figure, output_file, show, ColumnDataSource  
from bokeh.models import HoverTool  
from bokeh.io import output_notebook, show  
  
# this line tells jupyter notebook to put the plots in the notebook.  
%matplotlib inline  
  
# this line makes plots prettier on mac retina screens.  
%config InlineBackend.figure_format = 'retina'
```

0.1.8 3. Read and explore data; clean if necessary

```
In [3]: # read the csv file into a pandas dataframe  
golf = pd.read_csv("golf.csv")
```

```
In [4]: # get some basic info on the dataframe  
golf.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 194 entries, 0 to 193
Data columns (total 21 columns):
PLAYER NAME                194 non-null object
FEDEX CUP POINTS           194 non-null int64
AVERAGE DRIVING DISTANCE  194 non-null float64
Strokes Gained Off the Tee 194 non-null float64
Strokes gained tee to green 194 non-null float64
Strokes gained approach to the green 194 non-null float64
Strokes gained around the green 194 non-null float64
Strokes gained putting     194 non-null float64
Strokes gained total        194 non-null float64
BIRDIE TO BOGEY RATIO      194 non-null float64
birdie or better % from the rough 194 non-null float64
Scrambling average distance to the hole 194 non-null float64
World Ranking              194 non-null int64
Scoring Average            194 non-null float64
Percent of Available Purse Won 194 non-null float64
Official Money Won         194 non-null int64
Unnamed: 16                0 non-null float64
Unnamed: 17                0 non-null float64
Unnamed: 18                0 non-null float64
Unnamed: 19                0 non-null float64
Unnamed: 20                0 non-null float64
dtypes: float64(17), int64(3), object(1)
memory usage: 31.9+ KB

```

This is mostly encouraging. 194 rows and 21 columns of data.

Most are int and float which will be convenient for analysis.

The column names are easy to understand (at least for anyone who know the basics of golf!).

It's likely we have more columns of data then we will need to analyze the problem.

The first column PLAYER NAME and last five columns might be problematic.

Let's look at the head and tail to see if we can learn more.

```
In [5]: golf.head(5)
```

```

Out[5]:
   PLAYER NAME  FEDEX CUP POINTS  AVERAGE DRIVING DISTANCE  \
0  Aaron Baddeley                167                      291.6
1   Adam Hadwin                1198                      292.6
2   Adam Scott                 334                      304.3
3   Alex Cejka                 146                      280.6

```

| | | | |
|---|-----------------|-----|-------|
| 4 | Andres Gonzales | 119 | 297.3 |
|---|-----------------|-----|-------|

| | Strokes Gained Off the Tee | Strokes gained tee to green \ |
|---|----------------------------|-------------------------------|
| 0 | -0.618 | -0.184 |
| 1 | 0.223 | 0.856 |
| 2 | 0.356 | 0.772 |
| 3 | 0.124 | -0.416 |
| 4 | -0.526 | -0.483 |

| | Strokes gained approach to the green | Strokes gained around the green \ |
|---|--------------------------------------|-----------------------------------|
| 0 | 0.021 | 0.424 |
| 1 | 0.735 | 0.101 |
| 2 | 0.342 | 0.074 |
| 3 | -0.444 | 0.178 |
| 4 | 0.055 | -0.012 |

| | Strokes gained putting | Strokes gained total | BIRDIE TO BOGEY RATIO \ |
|---|------------------------|----------------------|-------------------------|
| 0 | -0.141 | -0.315 | 1.28 |
| 1 | 0.586 | 1.645 | 2.09 |
| 2 | 0.162 | 0.934 | 1.58 |
| 3 | -0.312 | -0.453 | 1.26 |
| 4 | 0.538 | 0.056 | 1.50 |

| | ... | Scrambling average distance to the hole | World Ranking \ |
|---|-----|---|-----------------|
| 0 | ... | 7.08 | 136 |
| 1 | ... | 8.25 | 46 |
| 2 | ... | 8.25 | 10 |
| 3 | ... | 6.50 | 186 |
| 4 | ... | 8.50 | 278 |

| | Scoring Average | Percent of Available Purse Won | Official Money Won \ |
|---|-----------------|--------------------------------|----------------------|
| 0 | 71.940 | 0.48 | 388322 |
| 1 | 69.948 | 2.95 | 2702158 |
| 2 | 70.001 | 1.57 | 891900 |
| 3 | 71.588 | 0.34 | 285056 |
| 4 | 71.283 | 0.28 | 204021 |

| | Unnamed: 16 | Unnamed: 17 | Unnamed: 18 | Unnamed: 19 | Unnamed: 20 |
|---|-------------|-------------|-------------|-------------|-------------|
| 0 | NaN | NaN | NaN | NaN | NaN |
| 1 | NaN | NaN | NaN | NaN | NaN |
| 2 | NaN | NaN | NaN | NaN | NaN |
| 3 | NaN | NaN | NaN | NaN | NaN |
| 4 | NaN | NaN | NaN | NaN | NaN |

[5 rows x 21 columns]

In [6]: golf.tail(5)

Out[6]:

| | PLAYER NAME | FEDEX CUP POINTS | AVERAGE DRIVING DISTANCE \ |
|-----|-------------------|------------------|----------------------------|
| 189 | Will MacKenzie | 78 | 279.8 |
| 190 | William McGirt | 432 | 285.5 |
| 191 | Xander Schauffele | 107 | 299.4 |
| 192 | Zac Blair | 225 | 272.1 |
| 193 | Zach Johnson | 258 | 286.7 |

| | Strokes Gained Off the Tee | Strokes gained tee to green \ |
|-----|----------------------------|-------------------------------|
| 189 | -0.213 | -0.794 |
| 190 | 0.433 | 0.290 |
| 191 | -0.005 | 0.183 |
| 192 | -0.075 | 0.171 |
| 193 | 0.051 | -0.009 |

| | Strokes gained approach to the green | Strokes gained around the green \ |
|-----|--------------------------------------|-----------------------------------|
| 189 | -0.408 | -0.173 |
| 190 | 0.192 | -0.168 |
| 191 | 0.233 | -0.045 |
| 192 | -0.121 | 0.344 |
| 193 | -0.446 | 0.386 |

| | Strokes gained putting | Strokes gained total | BIRDIE TO BOGEY RATIO \ |
|-----|------------------------|----------------------|-------------------------|
| 189 | -0.319 | -1.112 | 1.38 |
| 190 | -0.005 | 0.451 | 1.57 |
| 191 | 0.013 | 0.196 | 1.47 |
| 192 | 0.115 | 0.263 | 1.63 |
| 193 | 0.564 | 0.555 | 1.32 |

| | ... | Scrambling average distance to the hole | World Ranking \ |
|-----|-----|---|-----------------|
| 189 | ... | 8.92 | 439 |
| 190 | ... | 8.92 | 44 |
| 191 | ... | 8.83 | 337 |
| 192 | ... | 6.83 | 223 |
| 193 | ... | 6.75 | 56 |

| | Scoring Average | Percent of Available Purse Won | Official Money Won \ |
|-----|-----------------|--------------------------------|----------------------|
| 189 | 71.744 | 0.29 | 174280 |
| 190 | 70.788 | 1.02 | 969452 |
| 191 | 71.542 | 0.38 | 246710 |
| 192 | 70.931 | 0.47 | 475360 |
| 193 | 70.596 | 0.88 | 609429 |

| | Unnamed: 16 | Unnamed: 17 | Unnamed: 18 | Unnamed: 19 | Unnamed: 20 |
|-----|-------------|-------------|-------------|-------------|-------------|
| 189 | NaN | NaN | NaN | NaN | NaN |
| 190 | NaN | NaN | NaN | NaN | NaN |
| 191 | NaN | NaN | NaN | NaN | NaN |
| 192 | NaN | NaN | NaN | NaN | NaN |
| 193 | NaN | NaN | NaN | NaN | NaN |

[5 rows x 21 columns]

Ok, the data seems to be mostly clean.

The column names can be cleaned up to be more consistent format.

The last five columns are not needed.

So, let's delete those columns and do a final check on the data before moving forward.

```
In [7]: # rename columns for consistent format
# http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.rename.html
golf = golf.rename(columns = {'FEDEX CUP POINTS':"FedEx Cup Points",
                              'AVERAGE DRIVING DISTANCE':'Average Driving Distance',
                              'Strokes gained tee to green':'Strokes Gained Tee to Green',
                              'Strokes gained approach to the green':'Strokes Gained Approach to the Green',
                              'Strokes gained around the green':'Strokes Gained Around the Green',
                              'Strokes gained putting':'Strokes Gained Putting',
                              'Strokes gained total':'Strokes Gained Total',
                              'BIRDIE TO BOGEY RATIO':'Birdie to Bogey Ratio',
                              'birdie or better % from the rough':'Birdie or Better % From the Rough',
                              'Scrambling average distance to the hole':'Scrambling Average Distance to the Hole'})

# remove blank columns
# http://pandas.pydata.org/pandas-docs/version/0.17.0/generated/pandas.DataFrame.drop.html
golf = golf.drop('Unnamed: 16',axis=1)
golf = golf.drop('Unnamed: 17',axis=1)
golf = golf.drop('Unnamed: 18',axis=1)
golf = golf.drop('Unnamed: 19',axis=1)
golf = golf.drop('Unnamed: 20',axis=1)
```

```
In [8]: golf.head(2)
```

```
Out[8]:
```

| | PLAYER NAME | FedEx Cup Points | Average Driving Distance | \ |
|---|----------------|------------------|--------------------------|---|
| 0 | Aaron Baddeley | 167 | 291.6 | |
| 1 | Adam Hadwin | 1198 | 292.6 | |

| | Strokes Gained Off the Tee | Strokes Gained Tee to Green | \ |
|---|----------------------------|-----------------------------|---|
| 0 | -0.618 | -0.184 | |
| 1 | 0.223 | 0.856 | |

| | Strokes Gained Approach to the Green | Strokes Gained Around the Green | \ |
|---|--------------------------------------|---------------------------------|---|
| 0 | 0.021 | 0.424 | |
| 1 | 0.735 | 0.101 | |

| | Strokes Gained Putting | Strokes Gained Total | Birdie to Bogey Ratio | \ |
|--|------------------------|----------------------|-----------------------|---|
|--|------------------------|----------------------|-----------------------|---|

| | | | |
|---|--------|--------|------|
| 0 | -0.141 | -0.315 | 1.28 |
| 1 | 0.586 | 1.645 | 2.09 |

| | birdie or better % from the rough | Scrambling Average Distance to the Hole \ |
|---|-----------------------------------|---|
| 0 | 24.58 | 7.08 |
| 1 | 16.98 | 8.25 |

| | World Ranking | Scoring Average | Percent of Available Purse Won \ |
|---|---------------|-----------------|----------------------------------|
| 0 | 136 | 71.940 | 0.48 |
| 1 | 46 | 69.948 | 2.95 |

| | Official Money Won |
|---|--------------------|
| 0 | 388322 |
| 1 | 2702158 |

0.1.9 4. Analyze the data

The next line of code is simple but powerful. It gives us important insights about the data.

Observations about each column of data follow the data display.

```
In [9]: # start with key summary statistics
# http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.describe.html
golf.describe()
```

```
Out[9]:
```

| | FedEx Cup Points | Average Driving Distance | Strokes Gained Off the Tee \ |
|-------|------------------|--------------------------|------------------------------|
| count | 194.000000 | 194.000000 | 194.000000 |
| mean | 310.365979 | 290.468041 | 0.056495 |
| std | 320.255476 | 9.500581 | 0.435994 |
| min | 6.000000 | 268.300000 | -1.525000 |
| 25% | 103.250000 | 284.525000 | -0.205250 |
| 50% | 204.500000 | 290.550000 | 0.093000 |
| 75% | 408.250000 | 296.875000 | 0.347750 |
| max | 1903.000000 | 318.800000 | 1.342000 |

| | Strokes Gained Tee to Green | Strokes Gained Approach to the Green \ |
|-------|-----------------------------|--|
| count | 194.000000 | 194.000000 |
| mean | 0.120603 | 0.045392 |
| std | 0.799866 | 0.485720 |
| min | -2.631000 | -1.298000 |
| 25% | -0.384750 | -0.292000 |
| 50% | 0.126500 | 0.033000 |
| 75% | 0.567000 | 0.386750 |
| max | 2.830000 | 1.783000 |

| | Strokes Gained Around the Green | Strokes Gained Putting \ |
|-------|---------------------------------|--------------------------|
| count | 194.000000 | 194.000000 |
| mean | 0.019933 | -0.000907 |

| | | |
|-----|-----------|-----------|
| std | 0.274501 | 0.424551 |
| min | -0.818000 | -1.268000 |
| 25% | -0.162750 | -0.273750 |
| 50% | 0.022000 | 0.008500 |
| 75% | 0.198000 | 0.288000 |
| max | 0.622000 | 1.023000 |

| | Strokes Gained Total | Birdie to Bogey Ratio \ |
|-------|----------------------|-------------------------|
| count | 194.000000 | 194.000000 |
| mean | 0.120959 | 1.388660 |
| std | 0.900425 | 0.296322 |
| min | -3.271000 | 0.680000 |
| 25% | -0.425750 | 1.190000 |
| 50% | 0.079500 | 1.350000 |
| 75% | 0.650000 | 1.560000 |
| max | 2.953000 | 2.470000 |

| | birdie or better % from the rough \ |
|-------|-------------------------------------|
| count | 194.000000 |
| mean | 15.466753 |
| std | 4.395383 |
| min | 3.330000 |
| 25% | 12.797500 |
| 50% | 15.615000 |
| 75% | 18.360000 |
| max | 28.570000 |

| | Scrambling Average Distance to the Hole | World Ranking \ |
|-------|---|-----------------|
| count | 194.000000 | 194.000000 |
| mean | 8.144588 | 196.561856 |
| std | 1.023260 | 171.722148 |
| min | 5.170000 | 1.000000 |
| 25% | 7.520000 | 64.250000 |
| 50% | 8.250000 | 145.500000 |
| 75% | 8.830000 | 284.250000 |
| max | 10.830000 | 999.000000 |

| | Scoring Average | Percent of Available Purse Won | Official Money Won |
|-------|-----------------|--------------------------------|--------------------|
| count | 194.000000 | 194.000000 | 1.940000e+02 |
| mean | 71.220624 | 0.919124 | 7.309369e+05 |
| std | 0.853556 | 1.117250 | 8.245459e+05 |
| min | 68.868000 | 0.030000 | 1.552800e+04 |
| 25% | 70.634250 | 0.300000 | 2.143188e+05 |
| 50% | 71.234500 | 0.595000 | 4.588915e+05 |
| 75% | 71.749250 | 1.137500 | 9.367980e+05 |
| max | 74.531000 | 9.550000 | 5.346600e+06 |

0.1.10 Key findings from above:

FedEx Cup Points. This is for a year long competition and is a bit off the mark for the question at hand. Likely to eliminate.

Average Driving Distance. Wow! Pros hit the ball a long way. But, with a mean and median of 290 and std of 9.5, the competition is within a fairly tight range.

The next five columns look like the key ones to focus on for additional analysis. All measure gains or losses in strokes for different parts of the game. If we focus on std versus mean and median, we can see "Strokes Gained Tee to Green" has the most variance among the metrics. This is a tip off that it might be a key factor to understand further.

Strokes Gained Total and Scoring Average should be very highly correlated, and we will likely eliminate one after additional analysis.

The birdie and scrambling statistics have some wide ranges and deviations. Let's look at those more carefully.


Ah, the final two columns start to show us the Money! Let's see how they correlate with the rest.

0.1.11 Time to run a messy and too big pairplot and also a massive correlation matrix. It's gonna be a deluge of data charts, but it will help us to figure out what to keep and what to leave out going forward.

```
In [10]: # was getting a warning about font style that made display unattractive
# so, imported warning library at beginning of notebook
# then used this line of code so warning would not be displayed
# only added this after everything else completed
# thanks to Joe Kambourakis for help with the filterwarnings
# https://docs.python.org/3.1/library/warnings.html
warnings.filterwarnings("ignore")

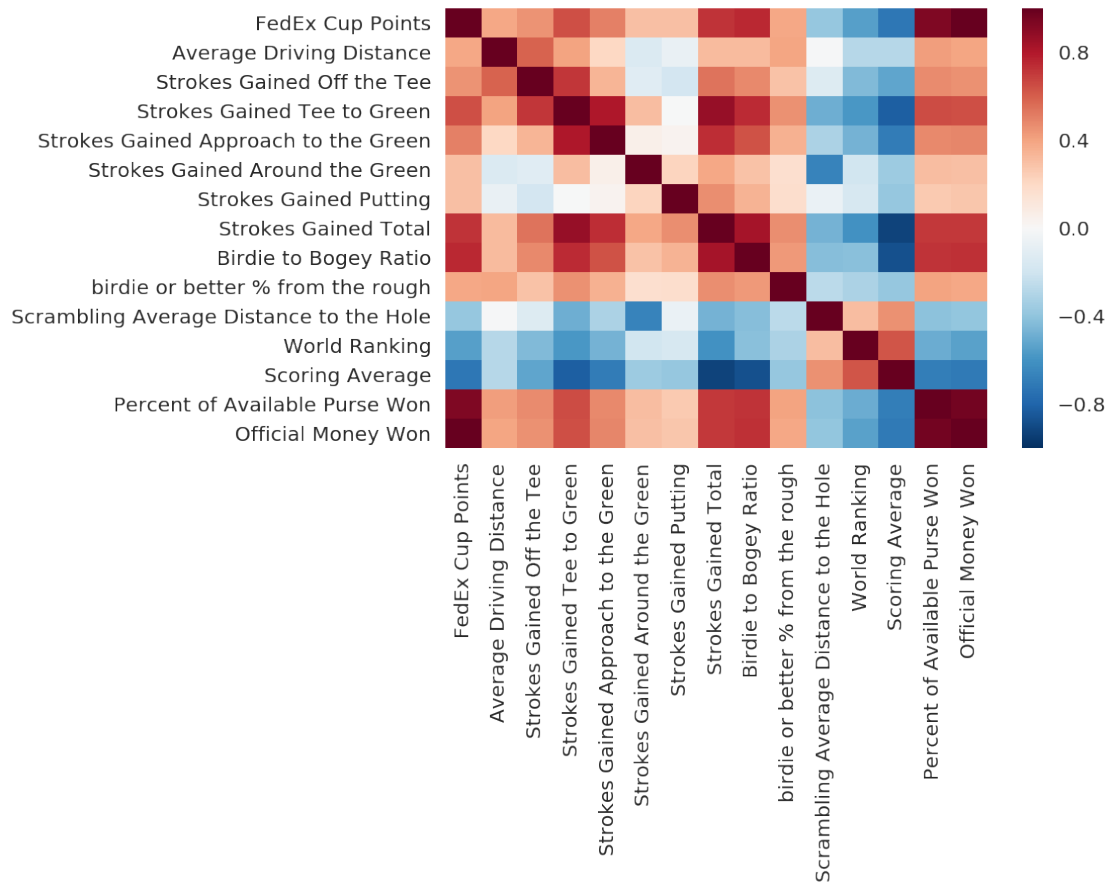
# this generates the pairplots
# http://seaborn.pydata.org/generated/seaborn.pairplot.html
sns.pairplot(golf)
```

```
Out[10]: <seaborn.axisgrid.PairGrid at 0x7f4b2c90ee48>
```



```
In [11]: # http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.corr.html
# http://seaborn.pydata.org/generated/seaborn.heatmap.html
golf_corr = golf.corr()
sns.heatmap(golf_corr)

Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4b2c2d9d30>
```



So, on the one hand, this is a gigantic mess. On the other, with a few simple lines of code, we are able to get tons of charts that help us to narrow the scope of the data to use going forward.

First, we see that FedEx Cup Points, Scoring Average, Percent of Available Purse Won and Official Money Won are all highly correlated. For the purposes of moving forward, let's just focus on the money...after all the project is focused on drive for show, putt for dough, so we really care about comparing strokes to Money won. We will create a new dataframe leaving the other proxies for money won out. Also, we will leave out World Ranking, which is a function of several years of factors, so it will be a messy variable to keep in.

Second, it seems like the real value will be comparing all of the Strokes Gained variables with Money Won. Those will be most relevant for the question at the beginning.

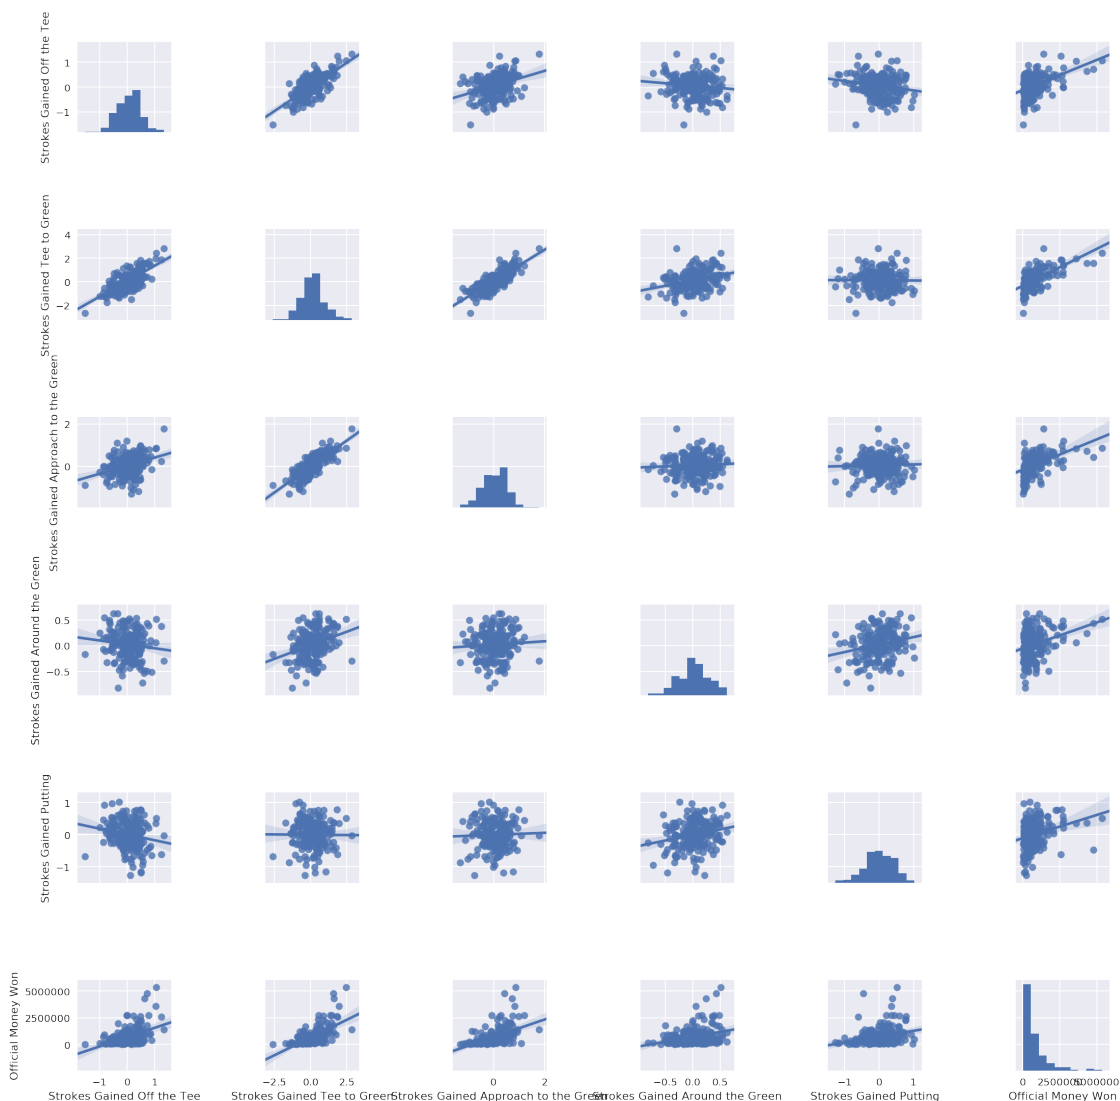
Third, if we focus on the bottom row of the correlation matrix as well as the bottom row of the pairplot output, we can see that Strokes Gained Tee to Green is looking like it is more highly correlated with Official Money Won than with the other Stokes Gained categories. Also, it looks like Strokes gained off the Tee is more highly correlated than Strokes Gained Putting.

Finally, Average Driving Distance is less correlated with Money Won than Strokes Gained off the Tee. So, we will drop that variable and just keep the Strokes Gained variables. Then, re-run the pairplots and correlation matrix.

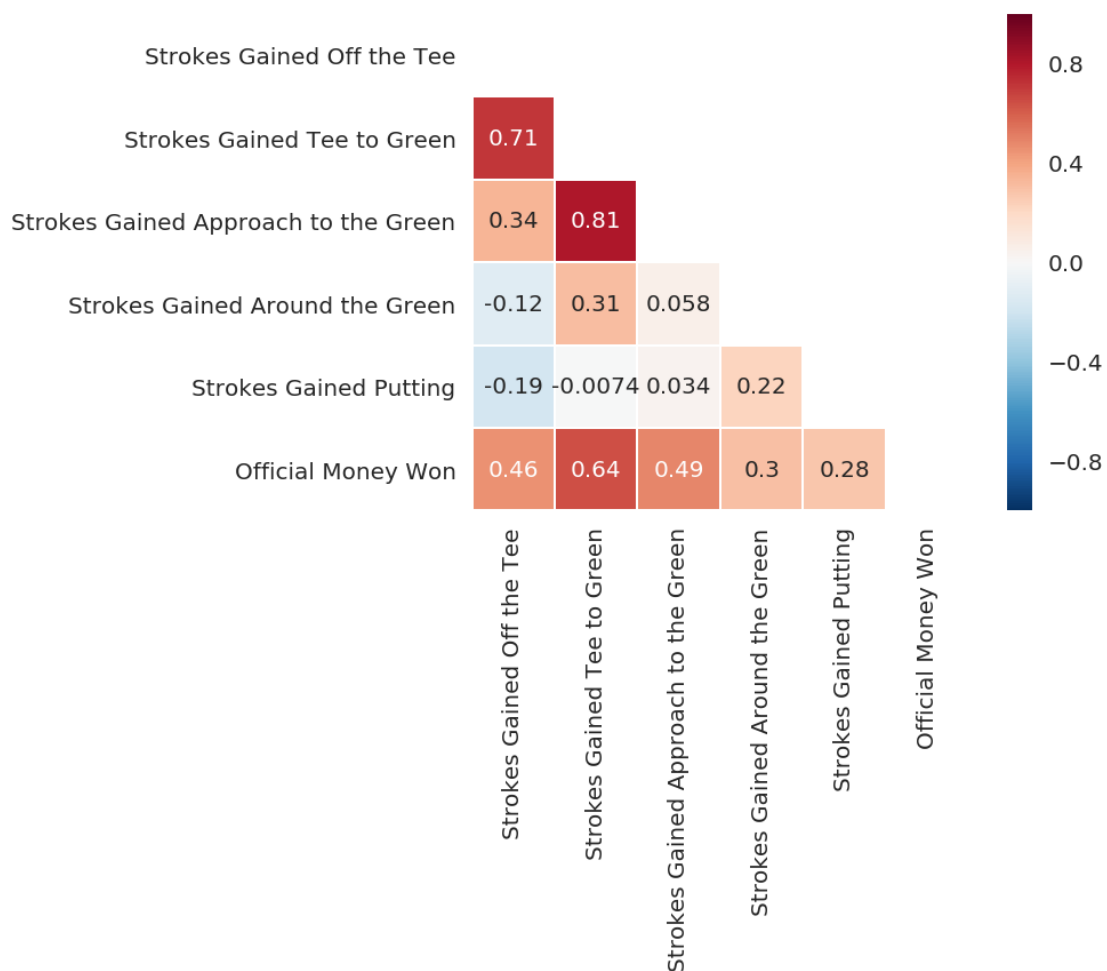
```
In [12]: golf = golf.drop('FedEx Cup Points',axis=1)
golf = golf.drop('Average Driving Distance',axis=1)
golf = golf.drop('Strokes Gained Total',axis=1)
golf = golf.drop('Birdie to Bogey Ratio',axis=1)
golf = golf.drop('birdie or better % from the rough',axis=1)
golf = golf.drop('Scrambling Average Distance to the Hole',axis=1)
golf = golf.drop('World Ranking',axis=1)
golf = golf.drop('Scoring Average',axis=1)
golf = golf.drop('Percent of Available Purse Won',axis=1)
```

```
In [13]: sns.pairplot(golf, kind='reg')
```

```
Out[13]: <seaborn.axisgrid.PairGrid at 0x7f4b1b69bcf8>
```



```
In [14]: # correlation matrix
# erik ellis (from data science bootcamp class)
# helped me get started on the concept of this layout
# I then spent a bunch of time with the documentation at:
# http://seaborn.pydata.org/generated/seaborn.heatmap.html
golf_corr = golf.corr()
mask = np.zeros_like(golf_corr)
mask[np.triu_indices_from(mask)] = True
with sns.axes_style("white"):
    ax = sns.heatmap(golf_corr, mask=mask, square=True, annot=True, linewidths=0.5)
```



- 0.1.12 Ah ha! Now, we are starting to get some key insights!
- 0.1.13 First, putting has the lowest correlation with Money Won at 0.28. So much for putt for dough!
- 0.1.14 Second, Strokes gained off the Tee (driving) is more correlated than putting at 0.46, but it's less correlated than Strokes gained tee to green at 0.64.
- 0.1.15 Third, Strokes gained approach to the green has a bit higher correlation (0.48) than off the tee (0.46). But, it's also highly correlated (0.81) with tee to green.
- 0.1.16 We seem to have the pieces of a story. Let's work on getting better visualizations to pull it together.
- 0.1.17 5. Create key data visualizations to tell a story
- 0.1.18 "Drive for Show, Putt for Dough" is a Myth

```
In [15]: # create a new DataFrame with just Money Won, Driving, Putting
        golf_chart = golf
        golf_chart = golf_chart.drop('Strokes Gained Approach to the Green',axis=1)
        golf_chart = golf_chart.drop('Strokes Gained Around the Green',axis=1)
        golf_chart = golf_chart.drop('Strokes Gained Tee to Green',axis=1)


        # dataframe with correlation data
        golf_chart_corr = golf_chart.corr()

In [16]: # output a simple bar chart of correlations with catchy title
        # http://pandas.pydata.org/pandas-docs/stable/visualization.html

        ax = golf_chart_corr['Official Money Won'].sort_values(ascending=True).plot.barh()

        ax.set_xlabel('Correlation Coefficients', fontsize=12)
        ax.set_title('"Drive for Show, Putt for Dough" is a Myth', fontsize=18)

        plt.show()
```



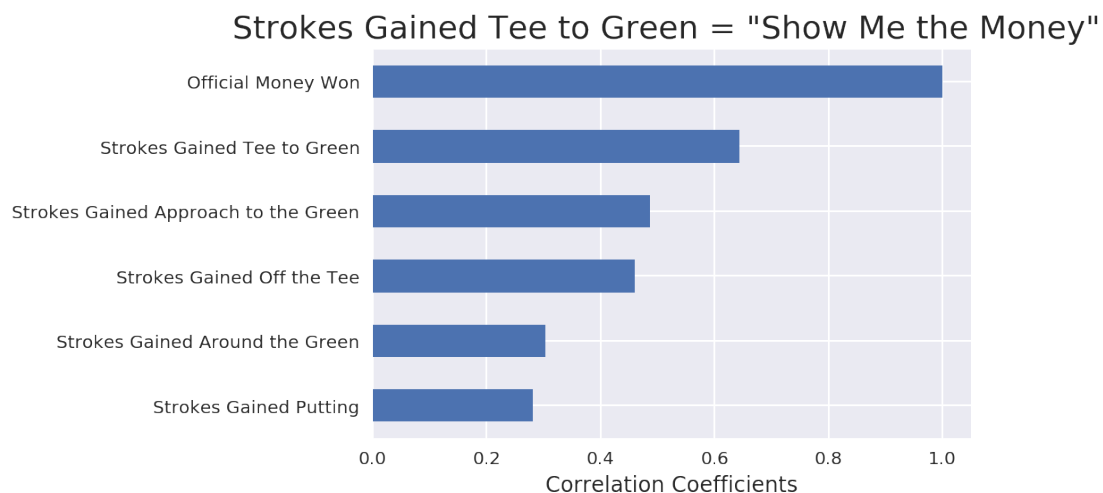
0.1.19 Strokes Gained Tee to Green = "Show Me the Money"

```
In [17]: # output a simple bar chart of correlations with catchy title
         # use the golf_corr dataframe from earlier that has all the Strokes Gained Variables

ax = golf_corr['Official Money Won'].sort_values(ascending=True).plot.barh()

ax.set_xlabel('Correlation Coefficients', fontsize=12)
ax.set_title('Strokes Gained Tee to Green = "Show Me the Money" ', fontsize=18)

plt.show()
```

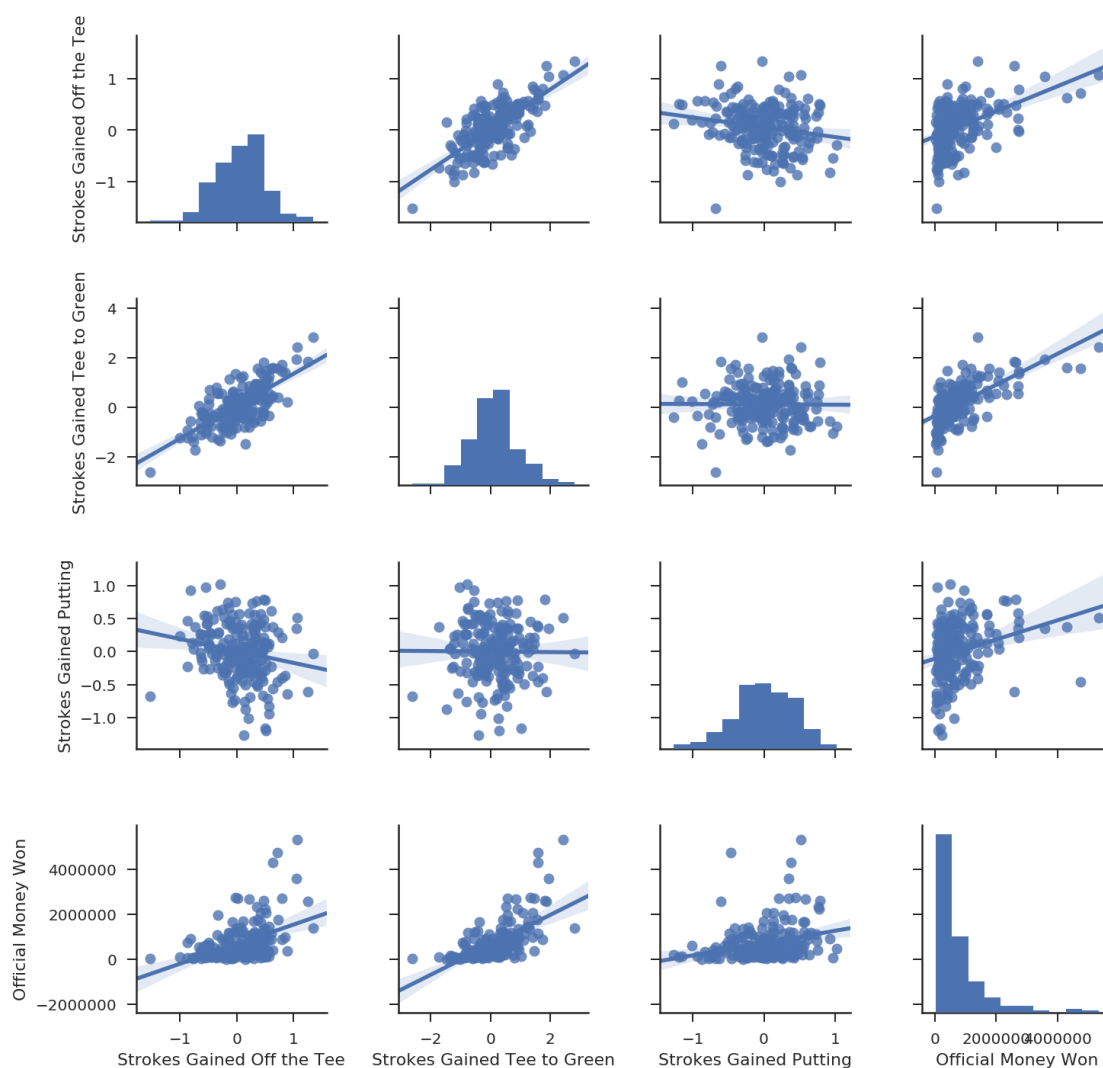


0.1.20 Pairplots allow us to visualize relationships between variables

```
In [18]: # create a new DataFrame with just Money Won, Driving, Putting, Tee to Green
golf_pairplot = golf
golf_pairplot = golf_pairplot.drop('Strokes Gained Approach to the Green',axis=1)
golf_pairplot = golf_pairplot.drop('Strokes Gained Around the Green',axis=1)
```

```
In [19]: sns.set(style="ticks", color_codes=True)
sns.pairplot(golf_pairplot, kind='reg')
```

```
Out[19]: <seaborn.axisgrid.PairGrid at 0x7f4b101acdd8>
```



0.1.21 Interactive bokeh charts are really cool!

0.1.22 On the chart below, highlight "box zoom" in the menu on the right.

0.1.23 Then use the cursor to highlight a box of the chart to zoom into.

```
In [20]: # this bokeh code allows us to box zoom and zero in on portions of the scatter plot
# we imported libraries at top of notebook; had we not, we would need the following here
# from bokeh.plotting import figure
# from bokeh.io import output_notebook, show
# http://bokeh.pydata.org/en/latest/docs/user_guide.html
# http://bokeh.pydata.org/en/latest/docs/user_guide/plotting.html

strokes = golf['Strokes Gained Tee to Green'] # x-values
money = golf['Official Money Won'] # y-values

# Set up the figure
p = figure(plot_width=600,
           plot_height=400,
           x_axis_label='Strokes Gained Tee to Green',
           y_axis_label='Official Money Won',
           title="Strokes Gained Tee to Green versus Official Money Won")

p.circle(strokes, money)

output_notebook()
show(p)
```

0.1.24 Cool chart, but the y-axis is hard to read.

0.1.25 Convert Official Money Won data to Thousand Dollars

```
In [21]: # convert Official Money Won to $ Thousands for better label display
golf['Official Money Won'] = golf["Official Money Won"].map(lambda x: int(x/1000))
golf['Strokes Gained Tee to Green'] = (golf['Strokes Gained Tee to Green']
                                       .map(lambda x: round(x,3)))
golf.head(2)
```

```
Out[21]:
```

| | PLAYER NAME | Strokes Gained Off the Tee | Strokes Gained Tee to Green \ |
|---|----------------|----------------------------|-------------------------------|
| 0 | Aaron Baddeley | -0.618 | -0.184 |
| 1 | Adam Hadwin | 0.223 | 0.856 |

| | Strokes Gained Approach to the Green | Strokes Gained Around the Green \ |
|---|--------------------------------------|-----------------------------------|
| 0 | 0.021 | 0.424 |
| 1 | 0.735 | 0.101 |

| | Strokes Gained Putting | Official Money Won |
|---|------------------------|--------------------|
| 0 | -0.141 | 388 |
| 1 | 0.586 | 2702 |

0.1.26 Now, use bokeh to create a hover tool.

0.1.27 Put cursor over any of the points in the scatter plot and learn player's name, money and tee to green score.

```
In [22]: # let's create a hover tool! very cool!
# we imported libraries at top of notebook; had we not, we would need the following here
# from bokeh.plotting import figure
# from bokeh.io import output_notebook, show
# from bokeh.plotting import figure, output_file, show, ColumnDataSource
# from bokeh.models import HoverTool
# http://bokeh.pydata.org/en/latest/docs/user_guide/tools.html#hover-tool
# output_file("toolbar.html")
# inspiration for this chart came from data science bootcamp

source = ColumnDataSource(
    data=dict(
        strokes = golf['Strokes Gained Tee to Green'],
        money = golf['Official Money Won'],
        name= golf['PLAYER NAME'],
    )
)

hover = HoverTool(
    tooltips=[
        ("name", "@name"),
        ("(strokes,money)", "@strokes, @money"),
    ]
)

p = figure(plot_width=600, plot_height=400, tools=[hover],
           title="Strokes Gained Tee to Green versus Official Money Won")

p.xaxis.axis_label = "Strokes Gained"
p.yaxis.axis_label = "Money Won (In $Thousand)"
p.xaxis.bounds = (-3,3)

p.circle('strokes', 'money', size=4, source=source)

output_notebook()

show(p)
```

0.1.28 Same chart as above, except x-axis is Strokes Gained Off the Tee.

```
In [23]: source = ColumnDataSource(
    data=dict(
        strokes = golf['Strokes Gained Off the Tee'],
        money = golf['Official Money Won'],
```

```

        name= golf['PLAYER NAME'],
    )
)

hover = HoverTool(
    tooltips=[
        ("name", "@name"),
        ("(strokes,money)", "@strokes, @money"),
    ]
)

p = figure(plot_width=550, plot_height=350, tools=[hover],
           title="Strokes Gained Off the Tee versus Money Won")

p.xaxis.axis_label = "Strokes Gained"
p.yaxis.axis_label = "Money Won (In $Thousand)"

p.circle('strokes', 'money', size=3, source=source)

output_notebook()

show(p)

```

0.1.29 Same chart as above, except x-axis is Strokes Gained Putting.

```

In [24]: source = ColumnDataSource(
        data=dict(
            strokes = golf['Strokes Gained Putting'],
            money = golf['Official Money Won'],
            name= golf['PLAYER NAME'],
        )
    )

hover = HoverTool(
    tooltips=[
        ("name", "@name"),
        ("(strokes,money)", "@strokes, @money"),
    ]
)

p = figure(plot_width=550, plot_height=350, tools=[hover],
           title="Strokes Gained Putting versus Money Won")

p.xaxis.axis_label = "Strokes Gained"
p.yaxis.axis_label = "Money Won (In $Thousand)"

p.circle('strokes', 'money', size=3, source=source)

```

```
output_notebook()
```

```
show(p)
```

0.1.30 It seems like the relationship between Strokes Gained Tee to Green and Official Money Won might be exponential.

0.1.31 Let's take the Log of Official Money Won and see if the correlations improve.

```
In [25]: # https://docs.scipy.org/doc/numpy-1.10.1/reference/generated/numpy.log.html
        golf['Log of Official Money Won'] = golf["Official Money Won"].apply(np.log)
```

```
In [26]: golf.head(2)
```

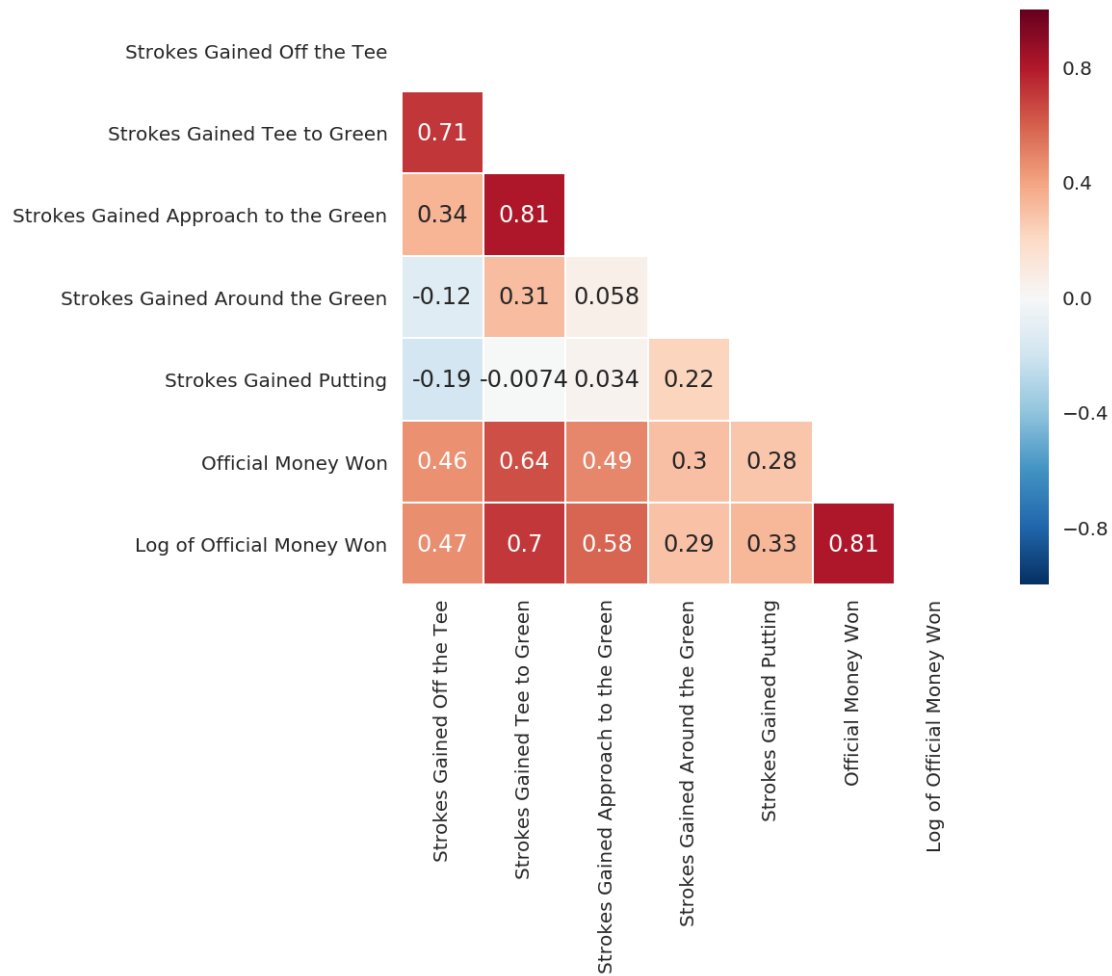
```
Out[26]:
```

| | PLAYER NAME | Strokes Gained Off the Tee | Strokes Gained Tee to Green \ |
|---|----------------|----------------------------|-------------------------------|
| 0 | Aaron Baddeley | -0.618 | -0.184 |
| 1 | Adam Hadwin | 0.223 | 0.856 |

| | Strokes Gained Approach to the Green | Strokes Gained Around the Green \ |
|---|--------------------------------------|-----------------------------------|
| 0 | 0.021 | 0.424 |
| 1 | 0.735 | 0.101 |

| | Strokes Gained Putting | Official Money Won | Log of Official Money Won |
|---|------------------------|--------------------|---------------------------|
| 0 | -0.141 | 388 | 5.961005 |
| 1 | 0.586 | 2702 | 7.901748 |

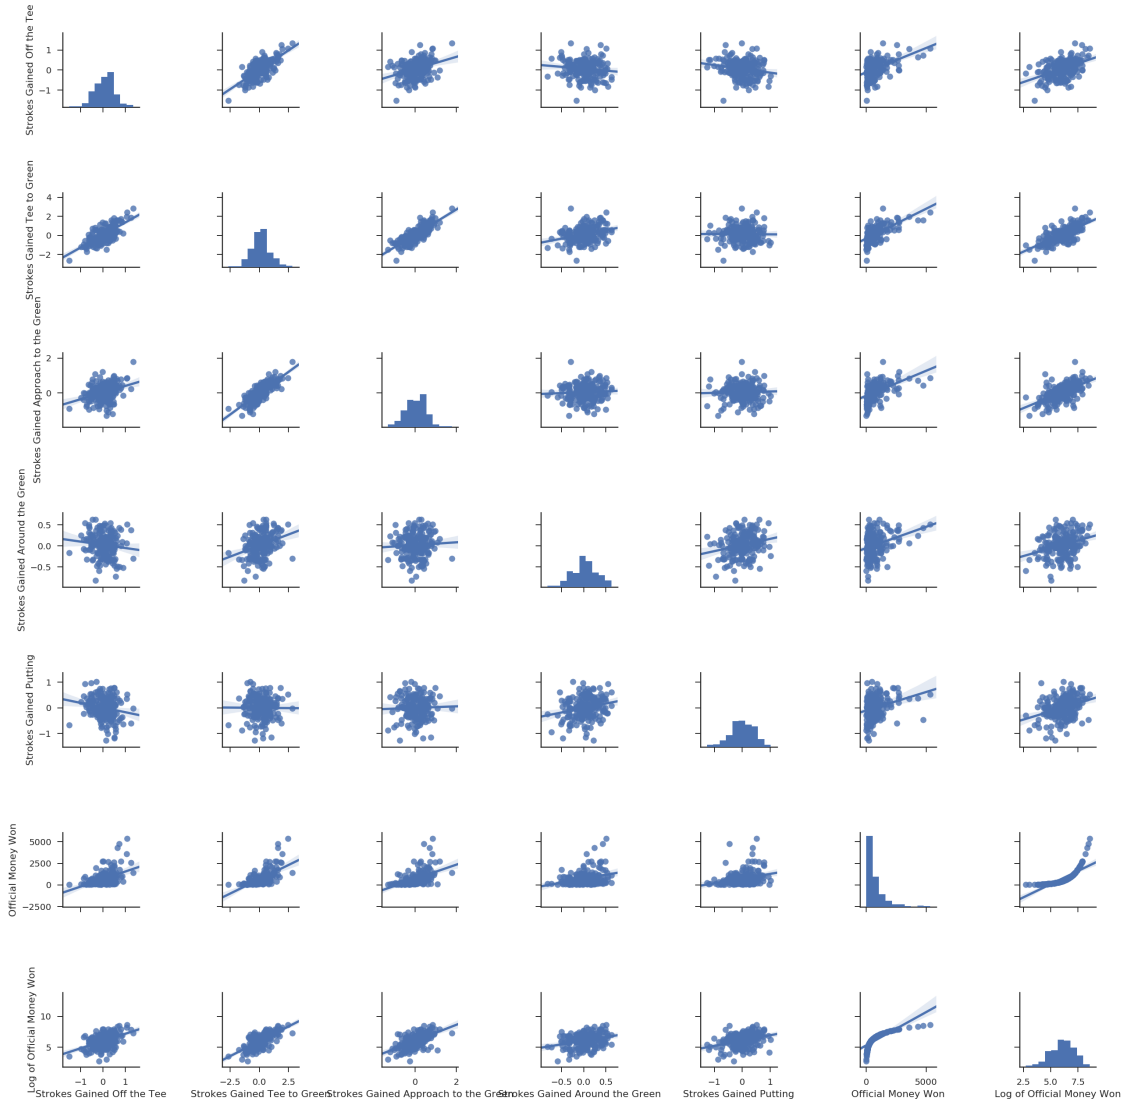
```
In [27]: # same code as earlier
        golf_corr = golf.corr()
        mask = np.zeros_like(golf_corr)
        mask[np.triu_indices_from(mask)] = True
        with sns.axes_style("white"):
            ax = sns.heatmap(golf_corr, mask=mask, square=True, annot=True, linewidths=0.5)
```



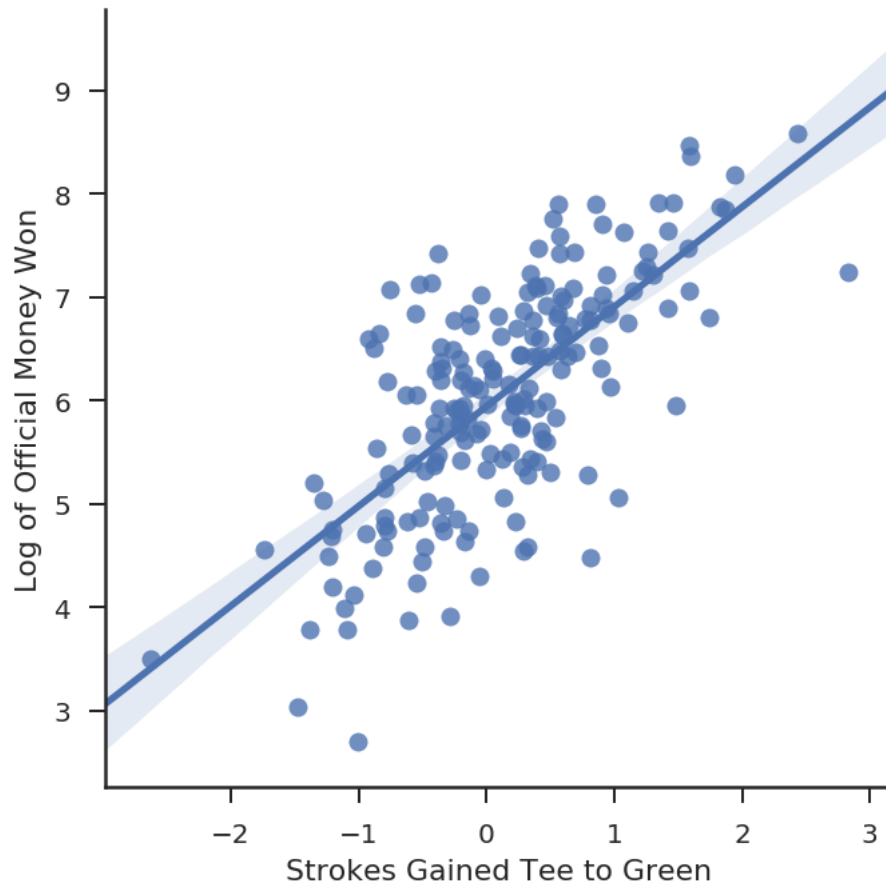
0.1.32 Looking at the last row above, we see Strokes Gained tee to Green has a 0.7 correlation with Log of Official Money Won! Now we are talking. Let's see the pairplots.

```
In [28]: sns.pairplot(golf, kind = 'reg')
```

```
Out[28]: <seaborn.axisgrid.PairGrid at 0x7f4b0a118940>
```



```
In [29]: # http://seaborn.pydata.org/generated/seaborn.lmplot.html
sns.lmplot(x='Strokes Gained Tee to Green', y='Log of Official Money Won', data=golf)
plt.show()
# this is really interesting because the lmplot allow us to
# It is intended as a convenient interface to fit regression models across conditional
# we are seeing a linear relationship now that we have log of official money won
```



0.2 Conclusions:

0.2.1 "Drive for Show, Putt for Dough" is a Myth

0.2.2 Strokes Gained Tee to Green = "Show Me the Money"

0.2.3 Pairplots Allow us to Visualize Relationships Between Variables

0.2.4 Bokeh makes really cool interactive visualizations

0.2.5 The visualizations lead us to look at the log of Official Money Won which has a 0.7 correlation with Strokes Gained Tee to Green. The last row of the pairplots gives a nice visual of the strong correlation.

0.2.6 Python and Golf are a match made in heaven

0.2.7 Thank you, Joe and Nenad, for a great course.

0.2.8 I learned a ton and had lots of fun!!!

In []: